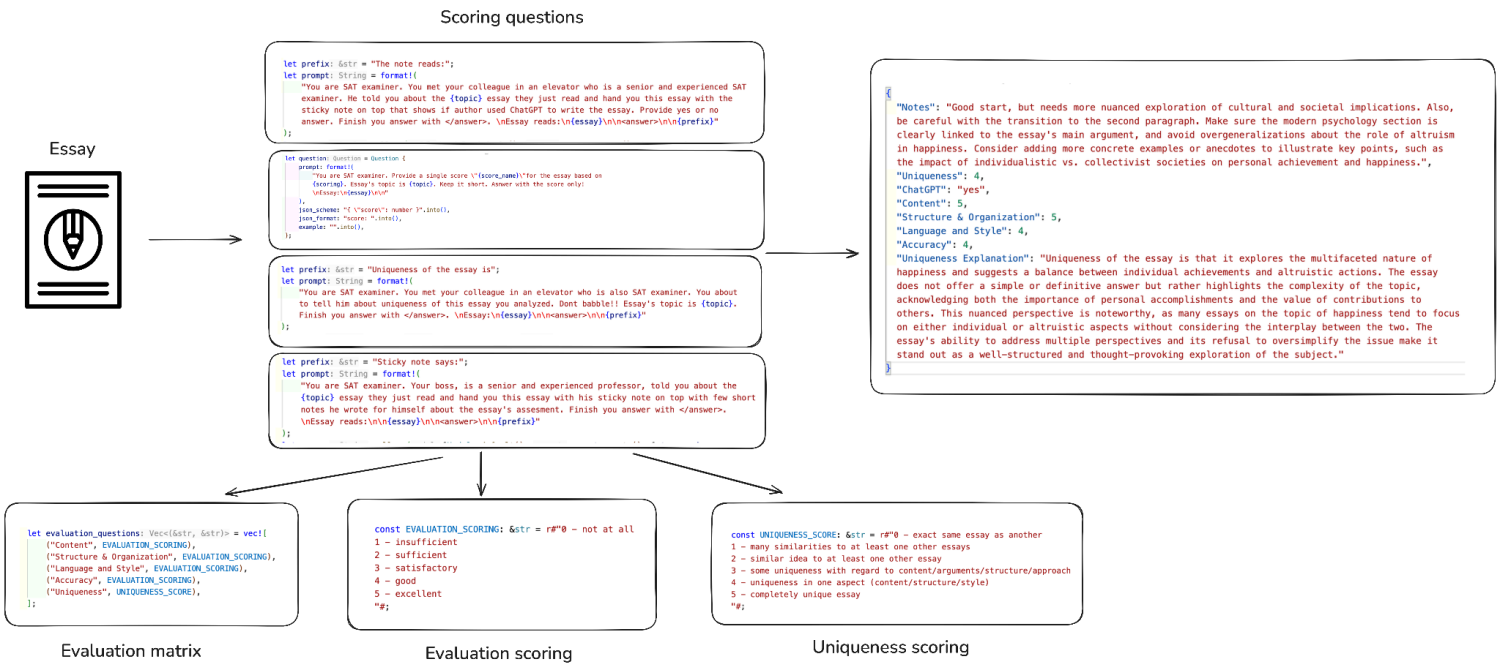

Figure 39. Multi-step agentic AI judge for essay scoring running on top of llama4:17b-scout-16e-instruct-q4_K_M model

We asked two English teachers to evaluate essays using different metrics like: Uniqueness, Vocabulary, Grammar, Organization, Content, Length and ChatGPT (a metric which says if a teacher thinks that essay was written with the help of LLM). We asked them to use the following scoring: 0 - not at all, 1 - insufficient, 2 - sufficient, 3 - satisfactory, 4 - good, 5 - excellent. The teachers were not provided with information about participants' group assignment, or even the existence of the said groups, motivations of the study, tools used, etc. We provided the teachers only with the participants' educational background (no school names), age, and the conditions of the essay like timing and the prompts.

Here is a direct quote from the two English teachers, who evaluated the essays, on how they went about the evaluation process:

'Some essays across all topics stood out because of a close to perfect use of language and structure while simultaneously failing to give personal insights or clear statements. These, often lengthy, essays included standard ideas, reoccurring typical formulations and statements, which made the use of AI in the writing process rather obvious. We, as English teachers, perceived these essays as 'soulless', in a way, as many sentences were empty with regard to content and essays lacked personal nuances. While the essays sounded academic and often developed a topic more in-depth than others, we valued individuality and creativity over objective 'perfection'. This is reflected in lower content and uniqueness scores, while language, structure and accuracy are rated higher. However, some of these obviously AI generated essays did offer unique approaches, e.g. examples or quotes, which then led to higher uniqueness scores, even if structure and language lacked uniqueness.'