We also created an AI judge using the same local LLM model, and asked it to evaluate essays in the same way human teachers did. We gave a system prompt that defined the AI judge as the writing expert. We found the AI judge was more statistically inclined to evaluate everything around a score of 4. See the distribution in Figure 40 below.

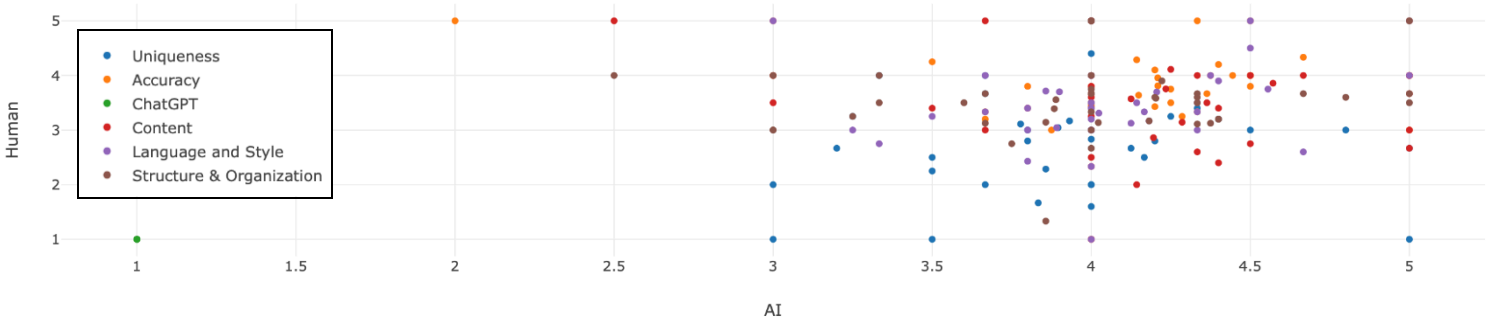

Figure 40. AI judge vs Human-Teacher Assessments Distribution. This scatter plot compares the average rankings given by human teachers and AI (LLM model) across different essay metrics. The X-axis represents the average scores assigned by the AI judge, while the Y-axis represents the average scores given by human teachers. Each dot on the plot corresponds to a specific essay metric, with the color of the dots differentiating between the metrics.

On average, human teachers assigned smaller scores to each metric except the ChatGPT metric, where teachers could not say exactly the LLMs were used to write the essays, however the AI judge assessed almost half of the essays as those that were written with the help of LLMs. See Figure 41 below.

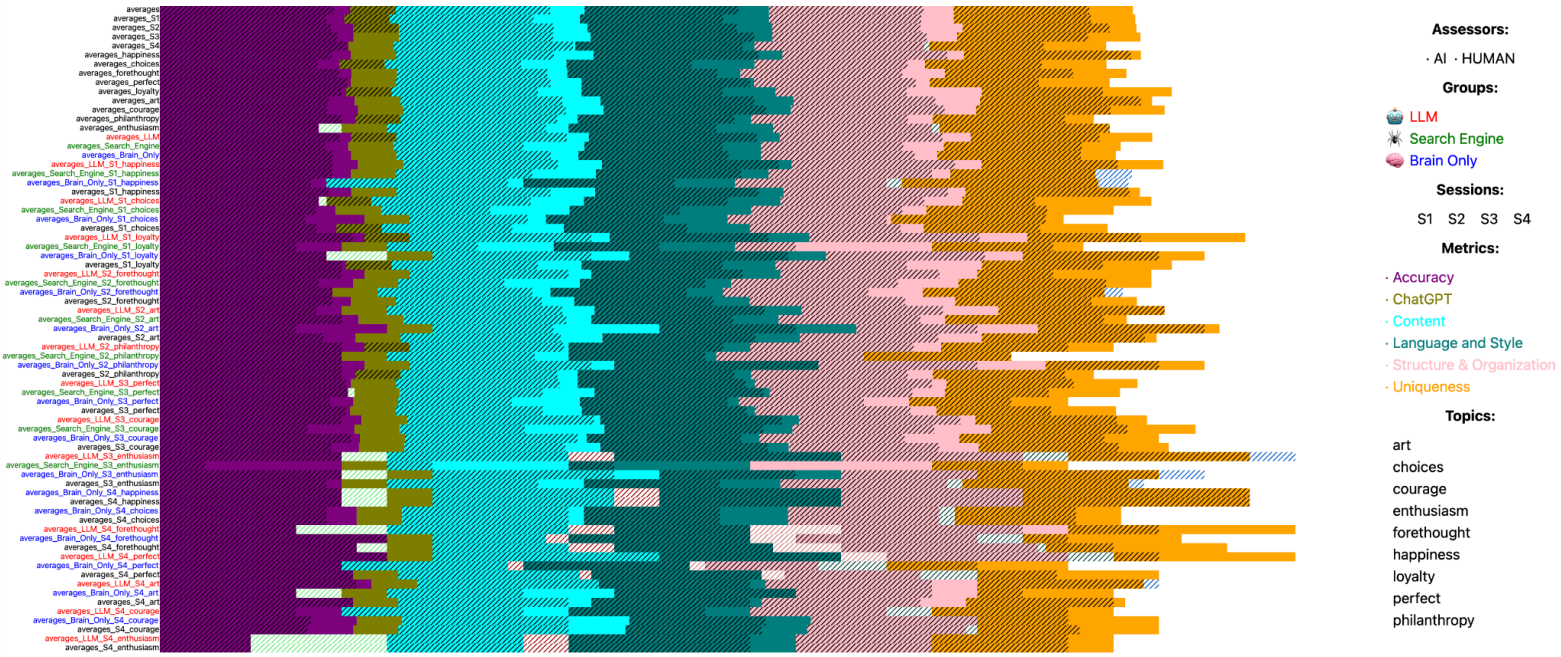

Figure 41. AI judge vs Human Teacher Assessments. This figure compares LLM-based AI assessments with human teacher evaluations for the essays across various metrics. The Y-axis shows the average scores assigned by each assessor, with the comparison highlighting consistency and discrepancies between AI and human judgments on the same set of essays. Solid color bars show AI judge assessments, while dashed overlaid bars show human-teacher