

The specific engineering of ChatGPT has made it quite compelling. But ultimately (at least until it can use outside tools) ChatGPT is “merely” pulling out some “coherent thread of text” from the “statistics of conventional wisdom” that it is accumulated. But it is amazing how human-like the results are. And as I have discussed, this suggests something that is at least scientifically very important: that human language (and the patterns of thinking behind it) are somehow simpler and more “law like” in their structure than we thought. ChatGPT has implicitly discovered it ¹.

This is not just a comment on the linguistic structure but also a deep philosophical reckoning on the nature of human intellect. Interactions with artificial minds can be unnerving not only because of the reasons listed in the opening of this essay. They are such also because they hold up a mirror to our own intellectual capacities, compelling us to question the bounds of our exceptionalism. We turned out to me much more like machines than we would care to admit. That is a concrete manifestation of liberatory alienation—the act of revealing an unpleasant truth about ourselves, causing loss, and yet ultimately liberating.

For eons, humans have been mightily impressed by their own intellectual and linguistic prowess (perhaps because we happen to be the only species with fully syntactic languages). We have often attributed these capabilities to divine influences, enshrining ourselves as the centerpiece of a cosmic masterpiece. This sense of self-importance has culminated in the belief that we were cast in the divine image, positioning ourselves on a unique pedestal. First, advancements in zoo psychology, and now, the neural networks puncture our inflated sense of self. Language generation, the crown jewel of human cognition, is now being replicated by AI-powered tools. This ability to mimic human language suggests that our linguistic prowess is not as enigmatic or complex as we once thought. Indeed, much of our communication is more patterned and predictable than we would like to admit, often reflecting a tendency to recycle and rephrase ideas and thoughts that we have heard or read before.

However, there is an enlightening side to this humbling realization. It helps us redefine our understanding of what it means to be human. Perhaps our essence, contrary to our previous grandiose self-perception, lies not in our linguistic abilities or intellectual prowess but rather in higher-level creative and discerning thinking and advanced ethical reasoning. These are qualities that, at least until now, have remained uniquely human, untouched by both the animal kingdom and the realm of AI. Machines help us notice and then discard what is machine-like in us.

4 Taking technology seriously

The process of technological advancement and its impact on human skills can be likened to the shedding of skin, a natural but unsettling transformation. What we once considered integral to our identity—akin to a limb—reveals itself to be more like a layer of dead skin, something that can be shed without compromising our core being. This shedding process is disconcerting because it forces us to confront existential questions about the nature of our true self. Is there a core that remains unshed, and if so, what does it consist of? Will we find this core to be as fulfilling and meaningful as we hope?

It is important to clarify that the claims made in this paper are not limited to narrow AI systems like ChatGPT but also encompass the potential impact of AGI. In Searle’s terms, the focus is not merely on weak AI, which simulates intelligent behavior, but also on the possibility of strong AI, which would possess genuine understanding and cognitive capabilities (Searle 1980). The implications of such systems would be far-reaching, potentially automating not only routine cognitive tasks but also more complex and creative endeavors. However, even in the absence of strong AI, the increasing sophistication of narrow AI systems across various domains could still bring about significant changes in the nature of work and human skills.

This paper argues that there is indeed a core, albeit smaller and more precious than we might have initially believed. However, this core is not something always already there, but a result of the gradual process of revelation that can also be understood as a process of construction. What is revealed/constructed is robust and possesses a beauty that transcends the layers we shed. As we continue to shed layers—whether they be manual skills or certain types of cognitive labor—we are not becoming less human but more human, defining our humanity against the tools we create. The shedding process, therefore, should not be viewed solely as a loss but as a form of liberation, a journey toward discovering the most authentic version of our collective self.

It is important to note that the argument here is not that higher capacities like creativity and agency cannot be technically automated. Rather, in the process of shedding mechanical layers and delegating tasks to AI, we are making the remaining human core increasingly pure, distilling the essence of what it means to be human. The tools may eventually develop human-like capabilities, but we will not necessarily need to delegate these core capacities to them. Instead, we will likely engage with art, music, and other creative pursuits produced by both humans and machines, appreciating the unique qualities of each. The small, irreducible human core is not subject to simple comparisons of

¹ Stephen Wolfram, “What Is ChatGPT Doing ... and Why Does It Work?” <https://writings.stephenwolfram.com/2023/02/what-is-chatgpt-doing-and-why-does-it-work/>